



[ANDROID](#) ▾
 [JAVA](#) ▾
 [JVM LANGUAGES](#) ▾
 [SOFTWARE DEVELOPMENT](#)
[AGILE](#)
[CAREER](#)
[COMMUNICATIONS](#)
[DEVOPS](#)
[META JCG](#) ▾

[Home](#) » [Java](#) » [Core Java](#) » 100 Multithreading and Concurrency Interview Questions and Answers – The ULTIMATE List (PDF Download)

ABOUT MARTIN MOIS



Martin is a Java EE enthusiast and works for an international operating company. He is interested in clean code and the software craftsmanship approach. He also strongly believes in automated testing and continuous integration.



100 Multithreading and Concurrency Interview Questions and Answers – The ULTIMATE List (PDF Download)

Posted by: Martin Mois
 in Core Java
 November 17th, 2014
 8 Comments
 6067 Views

Last Updated Jan. 26, 2019

EDITORIAL NOTE: *Concurrency is always a challenge for developers and writing concurrent programs can be extremely hard.*


There is a number of things that could potentially blow up and the complexity of systems rises considerably when concurrency is introduced.

However, the ability to write robust concurrent programs is a great tool in a developer's belt and can help build sophisticated, enterprise level applications.

In this article we will discuss different types of questions that can be used in a programming interview in order to assess a candidate's understanding of concurrency and multithreading.

The questions are not only Java specific, but revolve around general programming principles. Enjoy!

Programming Interview Coming Up?



Subscribe to our newsletter and download the **Ultimate** Multithreading and Concurrency interview questions and answers collection **right now!**

In order to get you prepared for your next Programming Interview, we have compiled a huge list of relevant Questions and their respective Answers. Besides studying them online you may download the eBook in PDF format!

Enter your e-mail...

☐ I agree to the Terms and Privacy Policy

Sign up

Table Of Contents

[Processes and Threads](#)
[Thread Objects](#)
[Defining and Starting a Thread](#)
[Pausing Execution with Sleep](#)
[Interrupts](#)
[Joins](#)
[Synchronization](#)
[Intrinsic Locks and Synchronization](#)
[Atomic Access](#)

NEWSLETTER

Insiders are already enjoying weekly up-to-date complimentary whitepapers!

Join them now to gain **exclusive** access to the latest news in the Java world, as well as insights about Android, Scala, and other related technologies.

Enter your e-mail...

☐ I agree to the Terms and Privacy Policy

Sign up

JOIN US



With **1,240,600** unique visitors and **500** authors placed among related sites and constantly being looked out for, you are encouraged to join. So if you have unique and interesting content then you should check out our **JCG** partners program. You can be a **guest writer** for Java Code Geek and improve your writing skills!

Processes and Threads

1. What do we understand by the term concurrency?

Concurrency is the ability of a program to execute several computations simultaneously. This can be achieved by distributing the computations over the available CPU cores of a machine or even over different machines within the same network.

[Learn more here:](#)

Java 8 Concurrency Tutorial

2. What is the difference between processes and threads?

A process is an execution environment provided by the operating system that has its own set of private resources (e.g. memory, open files, etc.). Threads, in contrast to processes, live within a process and share their resources (memory, open files, etc.) with the other threads of the process. The ability to share resources between different threads makes thread more suitable for tasks where performance is a significant requirement.

3. In Java, what is a process and a thread?

In Java, processes correspond to a running Java Virtual Machine (JVM) whereas threads live within the JVM and can be created and stopped by the Java application dynamically at runtime.

4. What is a scheduler?

A scheduler is the implementation of a scheduling algorithm that manages access of processes and threads to some limited resource like the processor or some I/O channel. The goal of most scheduling algorithms is to provide some kind of load balancing for the available processes/threads that guarantees that each process/thread gets an appropriate time frame to access the requested resource exclusively.

5. How many threads does a Java program have at least?

Each Java program is executed within the main thread; hence each Java application has at least one thread.

6. How can a Java application access the current thread?

The current thread can be accessed by calling the static method

```
currentThread()
```

of the JDK class

```
java.lang.Thread
```

:

```
1 public class MainThread {
2
3     public static void main(String[] args) {
4         long id = Thread.currentThread().getId();
5         String name = Thread.currentThread().getName();
6         ...
7     }
8 }
```

7. What properties does each Java thread have?

Each Java thread has the following properties:

- an identifier of type long that is unique within the JVM
- a name of type String
- a priority of type int
- a state of type

```
java.lang.Thread.State
```

- a thread group the thread belongs to

8. What is the purpose of thread groups?

Each thread belongs to a group of threads. The JDK class

```
java.lang.ThreadGroup
```



provides some methods to handle a whole group of Threads. With these methods we can, for example, interrupt all threads of a group or set their maximum priority.

9. What states can a thread have and what is the meaning of each state?

- NEW:

A thread that has not yet started is in this state.

- RUNNABLE:

A thread executing in the Java virtual machine is in this state.

- BLOCKED:

A thread that is blocked waiting for a monitor lock is in this state.

- WAITING:

A thread that is waiting indefinitely for another thread to perform a particular action is in this state.

- TIMED_WAITING:

A thread that is waiting for another thread to perform an action for up to a specified waiting time is in this state.

- TERMINATED:

A thread that has exited is in this state.

10. What are the benefits of multi-threaded programming?

Multithreading your code can help in the following areas:

- Improving Application Responsiveness
- Using Multiprocessors Efficiently
- Improving Program Structure
- Using Fewer System Resources

11. What are some common problems you have faced in multi-threading environment?

- Deadlock** – Two threads A and B, hold lock_A and lock_B respectively. Both of them want to access a resource R. In order to access R safely, both lock_A and lock_B are needed. But thread A needs lock_B and thread B needs lock_A. But both of them are not ready to give up the locks they hold. Hence there is no progress. It is deadlocked!
- Race conditions** – Consider the classic example of producer-consumer. What if you forget to lock before adding or removing an item from the queue? Imagine two threads A and B trying to add an item without locking. Thread A accesses the back of the queue. Scheduler then gives a chance to run thread B which successfully adds an item and updates the tail pointer. Now the tail pointer read by thread A is stale, but it thinks that it is the tail and adds the item. So the item added by B is LOST! The data structure is CORRUPTED! Worse, it might also result in memory leak while cleaning up.
- Data race** – Imagine a flag variable which should be set. Assume that you have put locks around to avoid race conditions. Now different threads want to set different values. Since the scheduler can schedule thread execution in any way, you don't know what the value of the flag is in the end.
- Starvation** – This is the problem caused by the thread scheduler. Some threads don't get a chance to run and complete or fail to obtain the required locks because other threads are given higher priority. They 'starve' for CPU cycles or other resources.
- Priority Inversion** – Imagine two threads A and B. A has higher priority than B and hence gets more CPU cycles than B. But while accessing a shared resource, B holds the lock which is also required by A and yields. Now A can't do anything without the lock and a lot of CPU cycles are wasted because B doesn't get enough cycles, but has the lock.

12. How do we set the priority of a thread?

The priority of a thread is set by using the method

```
setPriority(int)
```

. To set the priority to the maximum value, we use the constant

```
Thread.MAX_PRIORITY
```

and to set it to the minimum value we use the constant

```
Thread.MIN_PRIORITY
```

because these values can differ between different JVM implementations.

13. What is context-switching in multi-threading?



Pure multitasking doesn't exist. It is impossible to perform two mentally challenging tasks at the same time. Therefore, when we multitask, what we really do is constantly switch from one task to another. That's what context switching is.

x

The term originates from computer science. However, it is just as applicable to mental tasks performed by humans. After all, the human mind is similar to a CPU in many ways.

Just as a CPU running multi-threaded processes temporarily puts on hold an execution of a given thread while running another one, the human mind puts one task on hold in order to shift its focus to another one.

14. Difference between green thread and native thread in Java?

- Green threads refers to a model in which the Java virtual machine itself creates, manages, and context switches all Java threads within one operating system process. No operating system threads library is used.
- Native threads refers to a in which the Java virtual machine creates and manages Java threads using the operating system threads library – named libthread on UnixWare – and each Java thread is mapped to one threads library thread.

15. What do we understand by the term race condition?

A race condition describes constellations in which the outcome of some multi-threaded implementation depends on the exact timing behavior of the participating threads. In most cases it is not desirable to have such a kind of behavior, hence the term race condition also means that a bug due to missing thread synchronization leads to the differing outcome. A simple example for a race condition is the incrementation of an integer variable by two concurrent threads. As the operation consists of more than one single and atomic operation, it may happen that both threads read and increment the same value. After this concurrent incrementation the amount of the integer variable is not increased by two but only by one.

[Learn more here:](#)

java.util.concurrent.locks.Condition Example

16. What do you have to consider when passing object instances from one thread to another?

When passing objects between threads, you will have to pay attention that these objects are not manipulated by two threads at the same time. An example would be a

Map

implementation whose key/value pairs are modified by two concurrent threads. In order to avoid problems with concurrent modifications you can design an object to be immutable.

17. Is it possible to improve the performance of an application by the usage of multi-threading? Name some examples.

If we have more than one CPU core available, the performance of an application can be improved by multi-threading if it is possible to parallelize the computations over the available CPU cores. An example would be an application that should scale all images that are stored within a local directory structure. Instead of iterating over all images one after the other, a producer/consumer implementation can use a single thread to scan the directory structure and a bunch of worker threads that perform the actual scaling operation. Another example would be an application that mirrors some web page. Instead of loading one HTML page after the other, a producer thread can parse the first HTML page and issue the links it found into a queue. The worker threads monitor the queue and load the web pages found by the parser. While the worker threads wait for the page to get loaded completely, other threads can use the CPU to parse the already loaded pages and issue new requests.

18. What do we understand by the term scalability?

Scalability means the ability of a program to improve the performance by adding further resources to it.

19. Is it possible to compute the theoretical maximum speed up for an application by using multiple processors?

Amdahl's law provides a formula to compute the theoretical maximum speed up by providing multiple processors to an application. The theoretical speedup is computed by

$$S(n) = 1 / (B + (1-B)/n)$$

where

n

denotes the number of processors and

B

the fraction of the program that cannot be executed in parallel. When n converges against infinity, the term

$$(1-B)/n$$

converges against zero. Hence the formula can be reduced in this special case to

$$1/B$$

. As we can see, the theoretical maximum speedup behaves reciprocal to the fraction that has to be executed serially. This means the lower this fraction is, the more theoretical speedup can be achieved.

20. Provide an example why performance improvements for single-threaded applications can cause performance degradation for multi-threaded applications.

A prominent example for such optimizations is a

List

^
x

implementation that holds the number of elements as a separate variable. This improves the performance for single-threaded applications as the



```
size()
```

operation does not have to iterate over all elements but can return the current number of elements directly. Within a multi-threaded application the additional counter has to be guarded by a lock as multiple concurrent threads may insert elements into the list. This additional lock can cost performance when there are more updates to the list than invocations of the

```
size()
```

operation.

Thread Objects

Defining and Starting a Thread

21. How is a thread created in Java?

Basically, there are two ways to create a thread in Java.

The first one is to write a class that extends the JDK class

```
java.lang.Thread
```

and call its method

```
start()
```

:

```
01 public class MyThread extends Thread {
02
03     public MyThread(String name) {
04         super(name);
05     }
06
07     @Override
08     public void run() {
09         System.out.println("Executing thread "+Thread.currentThread().getName());
10     }
11
12     public static void main(String[] args) throws InterruptedException {
13         MyThread myThread = new MyThread("myThread");
14         myThread.start();
15     }
16 }
```

The second way is to implement the interface

```
java.lang.Runnable
```

and pass this implementation as a parameter to the constructor of

```
java.lang.Thread
```

:

```
01 public class MyRunnable implements Runnable {
02
03     public void run() {
04         System.out.println("Executing thread "+Thread.currentThread().getName());
05     }
06
07     public static void main(String[] args) throws InterruptedException {
08         Thread myThread = new Thread(new MyRunnable(), "myRunnable");
09         myThread.start();
10     }
11 }
```

22. Why should a thread not be stopped by calling its method

```
stop()
```

?

A thread should not be stopped by using the deprecated methods

```
stop()
```

of

```
java.lang.Thread
```

, as a call of this method causes the thread to unlock all monitors it has acquired. If any object protected by one of the released locks was in an inconsistent state, this state gets visible to all other threads. This can cause arbitrary behavior when other threads work this this inconsistent object.



23. Is it possible to start a thread twice?



No, after having started a thread by invoking its

```
start()
```

method, a second invocation of

```
start()
```

will throw an

```
IllegalThreadStateException
```

24. What is the output of the following code?

```
01 public class MultiThreading {
02
03     private static class MyThread extends Thread {
04
05         public MyThread(String name) {
06             super(name);
07         }
08
09         @Override
10         public void run() {
11             System.out.println(Thread.currentThread().getName());
12         }
13     }
14
15     public static void main(String[] args) {
16         MyThread myThread = new MyThread("myThread");
17         myThread.run();
18     }
19 }
```

The code above produces the output "main" and not "myThread". As can be seen in line two of the

```
main()
```

method, we invoke by mistake the method

```
run()
```

instead of

```
start()
```

. Hence, no new thread is started, but the method

```
run()
```

gets executed within the main thread.

25. What is a daemon thread?

A daemon thread is a thread whose execution state is not evaluated when the JVM decides if it should stop or not. The JVM stops when all user threads (in contrast to the daemon threads) are terminated. Hence daemon threads can be used to implement for example monitoring functionality as the thread is stopped by the JVM as soon as all user threads have stopped:

```
01 public class Example {
02
03     private static class MyDaemonThread extends Thread {
04
05         public MyDaemonThread() {
06             setDaemon(true);
07         }
08
09         @Override
10         public void run() {
11             while (true) {
12                 try {
13                     Thread.sleep(1);
14                 } catch (InterruptedException e) {
15                     e.printStackTrace();
16                 }
17             }
18         }
19     }
20
21     public static void main(String[] args) throws InterruptedException {
22         Thread thread = new MyDaemonThread();
23         thread.start();
24     }
25 }
```

The example application above terminates even though the daemon thread is still running in its endless while loop.

26. What is Java Thread Dump?

A Java thread dump is a way of finding out what every thread in the JVM is doing at a particular point in time.



This is especially useful if your Java application sometimes seems to hang when running under load, as an analysis of the dump will show where the threads are stuck.

x

You can generate a thread dump under Unix/Linux by running `kill -QUIT <pid>`, and under Windows by hitting `Ctrl + Break`.

27. Is it possible to convert a normal user thread into a daemon thread after it has been started?

A user thread cannot be converted into a daemon thread once it has been started. Invoking the method

```
thread.setDaemon(true)
```

on an already running thread instance causes a

```
IllegalThreadStateException
```

28. What do we understand by busy waiting?

Busy waiting means implementations that wait for an event by performing some active computations that let the thread/process occupy the processor although it could be removed from it by the scheduler. An example for busy waiting would be to spend the waiting time within a loop that determines the current time again and again until a certain point in time is reached:

```
01 Thread thread = new Thread(new Runnable() {
02     @Override
03     public void run() {
04         long millisToStop = System.currentTimeMillis() + 5000;
05         long currentTimeMillis = System.currentTimeMillis();
06         while (millisToStop > currentTimeMillis) {
07             currentTimeMillis = System.currentTimeMillis();
08         }
09     }
10 });
```

29. What are differences between wait() and sleep() method in Java?

wait():

- wait() method releases the lock.
- wait() is the method of Object class.
- wait() is the non-static method – public final void wait() throws InterruptedException { //... }
- wait() should be notified by notify() or notifyAll() methods.
- wait() method needs to be called from a loop in order to deal with false alarm.
- wait() method must be called from synchronized context (i.e. synchronized method or block), otherwise it will throw IllegalMonitorStateException

sleep():

- sleep() method doesn't release the lock.
- sleep() is the method of java.lang.Thread class.
- sleep() is the static method – public static void sleep(long millis, int nanos) throws InterruptedException { //... }
- after the specified amount of time, sleep() is completed.
- sleep() better not to call from loop (i.e. see code below).
- sleep() may be called from anywhere. there is no specific requirement.

30. What happens when an uncaught exception leaves the

```
run()
```

method?

I can happen that an unchecked exception escapes from the

```
run()
```

method. In this case the thread is stopped by the Java Virtual Machine. It is possible to catch this exception by registering an instance that implements the interface

```
UncaughtExceptionHandler
```

as an exception handler.

This is either done by invoking the static method

```
Thread.setDefaultUncaughtExceptionHandler(Thread.UncaughtExceptionHandler)
```

, which tells the JVM to use the provided handler in case there was no specific handler registered on the thread itself, or by invoking

```
setUncaughtExceptionHandler(Thread.UncaughtExceptionHandler)
```

on the thread instance itself.

31. What is the difference between the two interfaces

^
x

Runnable



and

Callable?

The interface

Runnable

defines the method

run()

without any return value whereas the interface

Callable

allows the method

call()

to return a value and to throw an exception.

32. What is a shutdown hook?

A shutdown hook is a thread that gets executed when the JVM shuts down. It can be registered by invoking

`addShutdownHook(Runnable)`

on the Runtime instance:

```
1 Runtime.getRuntime().addShutdownHook(new Thread() {  
2     @Override  
3     public void run() {  
4  
5     }  
6 });
```

Pausing Execution with Sleep

33. How can we prevent busy waiting?

One way to prevent busy waiting is to put the current thread to sleep for a given amount of time. This can be done by calling the method

`java.lang.Thread.sleep(long)`

by passing the number of milliseconds to sleep as an argument.

34. Can we use

`Thread.sleep()`

for real-time processing?

The number of milliseconds passed to an invocation of

`Thread.sleep(long)`

is only an indication for the scheduler how long the current thread does not need to be executed. It may happen that the scheduler lets the thread execute again a few milliseconds earlier or later depending on the actual implementation. Hence an invocation of

`Thread.sleep()`

should not be used for real-time processing.

35. What does the method

`Thread.yield()`

do?

An invocation of the static method

`Thread.yield()`

gives the scheduler a hint that the current thread is willing to free the processor. The scheduler is free to ignore this hint. As it is not defined which thread will get the processor after the invocation of

`Thread.yield()`

, it may even happen that the current thread becomes the “next” thread to be executed.

36. Which are use cases for the class

`java.util.concurrent.Future`



?



Instances of the class

java.util.concurrent.Future

are used to represent results of asynchronous computations whose result are not immediately available. Hence the class provides methods to check if the asynchronous computation has finished, canceling the task and to retrieve the actual result. The latter can be done with the two

get()

methods provided. The first

get()

methods takes no parameter and blocks until the result is available whereas the second

get()

method takes a timeout parameter that lets the method invocation return if the result does not get available within the given timeframe.

[Learn more here:](#)

java.util.concurrent.FutureTask Example

Java CompletionStage and CompletableFuture Example

37. How do we stop a thread in Java?

To stop a thread one can use a volatile reference pointing to the current thread that can be set to null by other threads to indicate the current thread should stop its execution:

```
01 private static class MyStopThread extends Thread {
02     private volatile Thread stopIndicator;
03
04     public void start() {
05         stopIndicator = new Thread(this);
06         stopIndicator.start();
07     }
08
09     public void stopThread() {
10         stopIndicator = null;
11     }
12
13     @Override
14     public void run() {
15         Thread thisThread = Thread.currentThread();
16         while(thisThread == stopIndicator) {
17             try {
18                 Thread.sleep(1000);
19             } catch (InterruptedException e) {
20             }
21         }
22     }
23 }
```

Interrupts

38. How can a thread be woken up that has been put to sleep before using

Thread.sleep()

?

The method

interrupt()

of

java.lang.Thread

interrupts a sleeping thread. The interrupted thread that has been put to sleep by calling

Thread.sleep()

is woken up by an

InterruptedException

:

```
01 public class InterruptExample implements Runnable {
02
03     public void run() {
04         try {
05             Thread.sleep(Long.MAX_VALUE);
06         } catch (InterruptedException e) {
07             System.out.println("[ "+Thread.currentThread().getName()+" ] Interrupted by exception!");
08         }
09     }
10 }
```



```

08     }
09 }
10
11 public static void main(String[] args) throws InterruptedException {
12     Thread myThread = new Thread(new InterruptExample(), "myThread");
13     myThread.start();
14
15     System.out.println "["+Thread.currentThread().getName()+" Sleeping in main thread for 5s..
16     Thread.sleep(5000);
17
18     System.out.println "["+Thread.currentThread().getName()+" Interrupting myThread");
19     myThread.interrupt();
20 }
21 }

```

39. How can a thread query if it has been interrupted?

If the thread is not within a method like

```
Thread.sleep()
```

that would throw an

```
InterruptedException
```

, the thread can query if it has been interrupted by calling either the static method

```
Thread.interrupted()
```

or the method

```
isInterrupted()
```

that it has inherited from

```
java.lang.Thread.
```

40. How should an

```
InterruptedException
```

be handled?

Methods like

```
sleep()
```

and

```
join()
```

throw an

```
InterruptedException
```

to tell the caller that another thread has interrupted this thread. In most cases this is done in order to tell the current thread to stop its current computations and to finish them unexpectedly. Hence ignoring the exception by catching it and only logging it to the console or some log file is often not the appropriate way to handle this kind of exception. The problem with this exception is, that the method

```
run()
```

of the Runnable interface does not allow that

```
run()
```

throws any exceptions. So just rethrowing it does not help. This means the implementation of

```
run()
```

has to handle this checked exception itself and this often leads to the fact that it is caught and ignored.

Joins

41. After having started a child thread, how do we wait in the parent thread for the termination of the child thread?

Waiting for a thread's termination is done by invoking the method

```
join()
```

on the thread's instance variable:

```

1 Thread thread = new Thread(new Runnable() {
2     @Override
3     public void run() {
4
5     }
6 });
7 thread.start();

```

```
8 | thread.join();
```

42. What is the output of the following program?

```
01 | public class MyThreads {
02 |
03 |     private static class MyDaemonThread extends Thread {
04 |
05 |         public MyDaemonThread() {
06 |             setDaemon(true);
07 |         }
08 |
09 |         @Override
10 |         public void run() {
11 |             try {
12 |                 Thread.sleep(1000);
13 |             } catch (InterruptedException e) {
14 |             }
15 |         }
16 |     }
17 |
18 |     public static void main(String[] args) throws InterruptedException {
19 |         Thread thread = new MyDaemonThread();
20 |         thread.start();
21 |         thread.join();
22 |         System.out.println(thread.isAlive());
23 |     }
24 | }
```

The output of the above code is "false". Although the instance of MyDaemonThread is a daemon thread, the invocation of

```
join()
```

causes the main thread to wait until the execution of the daemon thread has finished. Hence calling

```
isAlive()
```

on the thread instance reveals that the daemon thread is no longer running.

Synchronization

43. Differences between Synchronized method and Synchronized block?

- synchronized block reduce scope of lock, but synchronized method's scope of lock is whole method.
- synchronized block has better performance as only the critical section is locked but synchronized method has poor performance than block.
- synchronized block provide granular control over lock but synchronized method lock either on current object represented by this or class level lock.
- synchronized block can throw NullPointerException but synchronized method doesn't throw.
- synchronized block: synchronized(this){}
- synchronized method: public synchronized void fun(){}

44. What is a volatile keyword in Java and how is it different from the synchronized method in Java?

Using volatile force a thread to read and write variables directly from RAM memory. So when many thread are using the same volatile variable all them see the last version that is present in RAM memory and not a possible old copy in cache. When a thread enter a synchronized block it needs to take control of a monitor variable. All other threads wait until the first thread exit from the synchronized block. To ensure that all thread can see the same modifications all variables used in a synchronized block are read and written directly from the RAM memory instead from a cache copy.

45. For what purposes is the keyword synchronized used?

When you have to implement exclusive access to a resource, like some static value or some file reference, the code that works with the exclusive resource can be embraced with a synchronized block:

```
1 | synchronized (SynchronizedCounter.class) {
2 |     counter++;
3 | }
```

46. What is a semaphore?

A semaphore is a data structure that maintains a set of permits that have to be acquired by competing threads. Semaphores can therefore be used to control how many threads access a critical section or resource simultaneously. Hence the constructor of

```
java.util.concurrent.Semaphore
```

takes as first parameter the number of permits the threads compete about. Each invocation of its

```
acquire()
```

methods tries to obtain one of the available permits. The method

```
acquire()
```

without any parameter blocks until the next permit gets available. Later on, when the thread has finished its work on the critical resource, it can release the permit by invoking the method

```
release()
```



on an instance of Semaphore.



[Learn more here:](#)

Semaphores example limiting URL connections

java.util.concurrent.Semaphore Example

47. What is a

CountDownLatch

?

The SDK class

CountDownLatch

provides a synchronization aid that can be used to implement scenarios in which threads have to wait until some other threads have reached the same state such that all thread can start. This is done by providing a synchronized counter that is decremented until it reaches the value zero. Having reached zero the

CountDownLatch

instance lets all threads proceed. This can be either used to let all threads start at a given point in time by using the value 1 for the counter or to wait until a number of threads has finished. In the latter case the counter is initialized with the number of threads and each thread that has finished its work counts the latch down by one.

48. What is the difference between a

CountDownLatch

and a

CyclicBarrier

?

Both SDK classes maintain internally a counter that is decremented by different threads. The threads wait until the internal counter reaches the value zero and proceed from there on. But in contrast to the

CountDownLatch

the class

CyclicBarrier

resets the internal value back to the initial value once the value reaches zero. As the name indicates instances of

CyclicBarrier

can therefore be used to implement use cases where threads have to wait on each other again and again.

[Learn more here:](#)

Java CountDownLatch Example

java.util.concurrent.Phaser Example

Java.util.concurrent.CyclicBarrier Example

CountDownLatch example of a more general wait/notify mechanism

Intrinsic Locks and Synchronization

49. What intrinsic lock does a synchronized method acquire?

A synchronized method acquires the intrinsic lock for that method's object and releases it when the method returns. Even if the method throws an exception, the intrinsic lock is released. Hence a synchronized method is equal to the following code:

```
1 public void method() {  
2     synchronized(this) {  
3         ...  
4     }  
5 }
```

50. If two threads call a synchronized method on different object instances simultaneously, could one of these threads block?

Both methods lock the same monitor. Therefore, you can't simultaneously execute them on the same object from different threads (one of the two methods will block until the other is finished).

51. Can a constructor be synchronized?

No, a constructor cannot be synchronized. The reason why this leads to a syntax error is the fact that only the constructing thread should have access to the object being constructed.



52. What are Lock's benefits over synchronization?



The advantages of a lock are:

- it's possible to make them fair
- it's possible to make a thread responsive to interruption while waiting on a Lock object.
- it's possible to try to acquire the lock, but return immediately or after a timeout if the lock can't be acquired
- it's possible to acquire and release locks in different scopes, and in different orders

53. Can primitive values be used for intrinsic locks?

No, primitive values cannot be used for intrinsic locks.

54. Are intrinsic locks reentrant?

Yes, intrinsic locks can be accessed by the same thread again and again. Otherwise code that acquires a lock would have to pay attention that it does not accidentally tries to acquire a lock it has already acquired.

55. What do we understand by fair locks?

A fair lock takes the waiting time of the threads into account when choosing the next thread that passes the barrier to some exclusive resource. An example implementation of a fair lock is provided by the Java SDK:

```
java.util.concurrent.locks.ReentrantLock
```

. If the constructor with the Boolean flag set to true is used, the ReentrantLock grants access to the longest-waiting thread.

56. What kind of technique for reducing lock contention is used by the SDK class ReadWriteLock?

The SDK class

```
ReadWriteLock
```

uses the fact that concurrent threads do not have to acquire a lock when they want to read a value when no other thread tries to update the value. This is implemented by a pair of locks, one for read-only operations and one for writing operations. While the read-only lock may be obtained by more than one thread, the implementation guarantees that all read operation see an updated value once the write lock is released.

[Learn more here:](#)

Java ReentrantLock Example

Java ReadWriteLock Example

Java ReentrantReadWriteLock Example

Reentrant Lock example of a task runner

Reentrant ReadWriteLock example of value calculator

Atomic Access

57. What do we understand by an atomic operation?

An atomic operation is an operation that is either executed completely or not at all.

58. Is the statement `c++` atomic?

No, the incrementation of an integer variable consist of more than one operation. First we have to load the current value of `c`, increment it and then finally store the new value back. The current thread performing this incrementation may be interrupted in-between any of these three steps, hence this operation is not atomic.

59. What operations are atomic in Java?

The Java language provides some basic operations that are atomic and that therefore can be used to make sure that concurrent threads always see the same value:

- Read and write operations to reference variables and primitive variables (except long and double)
- Read and write operations for all variables declared as volatile

Liveness

Deadlock

60. What do we understand by a deadlock?

A deadlock is a situation in which two (or more) threads are each waiting on the other thread to free a resource that it has locked, while the thread itself has locked a resource the other thread is waiting on:



- Thread 1: locks resource A, waits for resource B
- Thread 2: locks resource B, waits for resource A



61. What are the requirements for a deadlock situation?

In general the following requirements for a deadlock can be identified:

- Mutual exclusion: There is a resource which can be accessed only by one thread at any point in time.
- Resource holding: While having locked one resource, the thread tries to acquire another lock on some other exclusive resource.
- No preemption: There is no mechanism, which frees the resource if one thread holds the lock for a specific period of time.
- Circular wait: During runtime a constellation occurs in which two (or more) threads are each waiting on the other thread to free a resource that it has locked.

62. Is it possible to prevent deadlocks at all?

In order to prevent deadlocks one (or more) of the requirements for a deadlock has to be eliminated:

- Mutual exclusion: In some situation it is possible to prevent mutual exclusion by using optimistic locking.
- Resource holding: A thread may release all its exclusive locks, when it does not succeed in obtaining all exclusive locks.
- No preemption: Using a timeout for an exclusive lock frees the lock after a given amount of time.
- Circular wait: When all exclusive locks are obtained by all threads in the same sequence, no circular wait occurs.

63. Is it possible to implement a deadlock detection?

When all exclusive locks are monitored and modelled as a directed graph, a deadlock detection system can search for two threads that are each waiting on the other thread to free a resource that it has locked. The waiting threads can then be forced by some kind of exception to release the lock the other thread is waiting on.

Starvation and Livelock

64. What is a livelock?

A livelock is a situation in which two or more threads block each other by responding to an action that is caused by another thread. In contrast to a deadlock situation, where two or more threads wait in one specific state, the threads that participate in a livelock change their state in a way that prevents progress on their regular work. An example would be a situation in which two threads try to acquire two locks, but release a lock they have acquired when they cannot acquire the second lock. It may now happen that both threads concurrently try to acquire the first thread. As only one thread succeeds, the second thread may succeed in acquiring the second lock. Now both threads hold two different locks, but as both want to have both locks, they release their lock and try again from the beginning. This situation may now happen again and again.

65. What do we understand by thread starvation?

Threads with lower priority get less time for execution than threads with higher priority. When the threads with lower priority performs a long enduring computations, it may happen that these threads do not get enough time to finish their computations just in time. They seem to "starve" away as threads with higher priority steal them their computation time.

66. Can a synchronized block cause thread starvation?

The order in which threads can enter a synchronized block is not defined. So in theory it may happen that in case many threads are waiting for the entrance to a synchronized block, some threads have to wait longer than other threads. Hence they do not get enough computation time to finish their work in time.

Guarded Blocks

67. Which two methods that each object inherits from

```
java.lang.Object
```

can be used to implement a simple producer/consumer scenario?

When a worker thread has finished its current task and the queue for new tasks is empty, it can free the processor by acquiring an intrinsic lock on the queue object and by calling the method

```
wait()
```

. The thread will be woken up by some producer thread that has put a new task into the queue and that again acquires the same intrinsic lock on the queue object and calls

```
notify()
```

on it.

68. What is the difference between

```
notify()
```

and

```
notifyAll()
```



?



Both methods are used to wake up one or more threads that have put themselves to sleep by calling

```
wait()
```

. While

```
notify()
```

only wakes up one of the waiting threads,

```
notifyAll()
```

wakes up all waiting threads.

69. How it is determined which thread wakes up by calling

```
notify()
```

?

It is not specified which threads will be woken up by calling

```
notify()
```

if more than one thread is waiting. Hence code should not rely on any concrete JVM implementation.

70. Is the following code that retrieves an integer value from some queue implementation correct?

```
01 public Integer getNextInt() {
02     Integer retVal = null;
03     synchronized (queue) {
04         try {
05             while (queue.isEmpty()) {
06                 queue.wait();
07             }
08         } catch (InterruptedException e) {
09             e.printStackTrace();
10         }
11     }
12     synchronized (queue) {
13         retVal = queue.poll();
14         if (retVal == null) {
15             System.err.println("retVal is null");
16             throw new IllegalStateException();
17         }
18     }
19     return retVal;
20 }
```

Although the code above uses the queue as object monitor, it does not behave correctly in a multi-threaded environment. The reason for this is that it has two separate synchronized blocks. When two threads are woken up in line 6 by another thread that calls

```
notifyAll()
```

, both threads enter one after the other the second synchronized block. In this second block the queue has now only one new value, hence the second thread will poll on an empty queue and get null as return value.

Immutable Objects

71. What is an Immutable Object

- Immutable objects can be published through any mechanism
- Immutable objects can be used safely by any thread without additional synchronization, even when synchronization is not used to publish them.

72. How to Create an Immutable Object

To create an immutable object you need to:

- Don't add any setter method
- Declare all fields final and private
- If a field is a mutable object create defensive copies of it for getter methods
- If a mutable object passed to the constructor must be assigned to a field create a defensive copy of it
- Don't allow sub-classes to override methods.

73. Which rules do you have to follow in order to implement an immutable class?

- All fields should be final and private.
- There should be no setter methods.
- The class itself should be declared final in order to prevent subclasses to violate the principle of immutability.
- If fields are not of a primitive type but a reference to another object:



- There should not be a getter method that exposes the reference directly to the caller.
- Don't change the referenced objects (or at least changing these references is not visible to clients of the object).



Lock Objects

74. Is it possible to check whether a thread holds a monitor lock on some given object?

The class

```
java.lang.Thread
```

provides the static method

```
Thread.holdsLock(Object)
```

that returns true if and only if the current thread holds the lock on the object given as argument to the method invocation.

75. What do we understand by lock contention?

Lock contention occurs, when two or more threads are competing in the acquisition of a lock. The scheduler has to decide whether it lets the thread, which has to wait sleeping and performs a context switch to let another thread occupy the CPU, or if letting the waiting thread busy-waiting is more efficient. Both ways introduce idle time to the inferior thread.

76. Which techniques help to reduce lock contention?

In some cases lock contention can be reduced by applying one of the following techniques:

- The scope of the lock is reduced.
- The number of times a certain lock is acquired is reduced (lock splitting).
- Using hardware supported optimistic locking operations instead of synchronization.
- Avoid synchronization where possible.
- Avoid object pooling.

77. Which technique to reduce lock contention can be applied to the following code?

```
1 synchronized (map) {  
2     UUID.randomUUID() = UUID.randomUUID();  
3     Integer value = Integer.valueOf(42);  
4     String key =.randomUUID.toString();  
5     map.put(key, value);  
6 }
```

The code above performs the computation of the random UUID and the conversion of the literal 42 into an Integer object within the synchronized block, although these two lines of code are local to the current thread and do not affect other threads. Hence they can be moved out of the synchronized block:

```
1 UUID.randomUUID() = UUID.randomUUID();  
2 Integer value = Integer.valueOf(42);  
3 String key =.randomUUID.toString();  
4 synchronized (map) {  
5     map.put(key, value);  
6 }
```

78. Explain by an example the technique lock splitting.

Lock splitting may be a way to reduce lock contention when one lock is used to synchronize access to different aspects of the same application. Suppose we have a class that implements the computation of some statistical data of our application. A first version of this class uses the keyword synchronized in each method signature in order to guard the internal state before corruption by multiple concurrent threads. This also means that each method invocation may cause lock contention as other threads may try to acquire the same lock simultaneously. But it may be possible to split the lock on the object instance into a few smaller locks for each type of statistical data within each method. Hence thread T1 that tries to increment the statistical data D1 does not have to wait for the lock while thread T2 simultaneously updates the data D2.

79. What do we understand by lock striping?

In contrast to lock splitting, where we introduce different locks for different aspects of the application, lock striping uses multiple locks to guard different parts of the same data structure. An example for this technique is the class

```
ConcurrentHashMap
```

from JDK's

```
java.util.concurrent
```

package. The

```
Map
```

implementation uses internally different buckets to store its values. The bucket is chosen by the value's key.

```
ConcurrentHashMap
```



now uses different locks to guard different hash buckets. Hence one thread that tries to access the first hash bucket can acquire the lock for this bucket, while another thread can simultaneously access a second bucket. In contrast to a synchronized version of

x

HashMap

this technique can increase the performance when different threads work on different buckets.

80. What is Lock interface in Java Concurrency API?

A `java.util.concurrent.locks.Lock` interface is used to as a thread synchronization mechanism similar to synchronized blocks. New Locking mechanism is more flexible and provides more options than a synchronized block.

Main differences between a Lock and a synchronized block are following:

- Guarantee of sequence – Synchronized block does not provide any guarantee of sequence in which waiting thread will be given access. Lock interface handles it.
- No timeout – Synchronized block has no option of timeout if lock is not granted. Lock interface provides such option.
- Single method – Synchronized block must be fully contained within a single method whereas a lock interface's methods `lock()` and `unlock()` can be called in different methods.

Executors

Executor Interfaces

81. Pros of ExecutorService over Timer

- Timer can't take advantage of available CPU cores unlike ExecutorService especially with multiple tasks using flavours of ExecutorService like `ForkJoinPool`
- ExecutorService provides collaborative API if you need coordination between multiple tasks. Assume that you have to submit N number of worker tasks and wait for completion of all of them. You can easily achieve it with `invokeAll` API. If you want to achieve the same with multiple Timer tasks, it would be not simple.
- `ThreadPoolExecutor` provides better API for management of Thread life cycle.

82. What is the relation between the two interfaces Executor and ExecutorService?

The interface

Executor

only defines one method:

`execute(Runnable)`

. Implementations of this interface will have to execute the given Runnable instance at some time in the future. The

ExecutorService

interface is an extension of the

Executor

interface and provides additional methods to shut down the underlying implementation, to await the termination of all submitted tasks and it allows submitting instances of

Callable

.

83. What happens when you

`submit()`

a new task to an ExecutorService instance whose queue is already full?

As the method signature of

`submit()`

indicates, the

ExecutorService

implementation is supposed to throw a

`RejectedExecutionException`

.

84. What is a ScheduledExecutorService?

^ x

The interface

ScheduledExecutorService



extends the interface

ExecutorService

and adds method that allow to submit new tasks to the underlying implementation that should be executed a given point in time. There are two methods to schedule one-shot tasks and two methods to create and execute periodic tasks.

[Learn more here:](#)

Java CompletionService Example

Java ExecutorService Example – Tutorial

Java ScheduledExecutorService Example

Java RunnableScheduledFuture Example

java.util.concurrent.ThreadFactory Example

java.util.concurrent.RejectedExecutionException – How to solve RejectedExecutionException

Exchanger example passing logs to a background logger

java.util.concurrent.RejectedExecutionHandler Example

java.util.concurrent.ScheduledThreadPoolExecutor Example

Thread Pools

85. Do you know an easy way to construct a thread pool with 5 threads that executes some tasks that return a value?

The SDK provides a factory and utility class

Executors

whose static method

newFixedThreadPool(int nThreads)

allows the creation of a thread pool with a fixed number of threads (the implementation of

MyCallable

is omitted):

```
01 public static void main(String[] args) throws InterruptedException, ExecutionException {
02     ExecutorService executorService = Executors.newFixedThreadPool(5);
03     Future<Integer>[] futures = new Future[5];
04     for (int i = 0; i < futures.length; i++) {
05         futures[i] = executorService.submit(new MyCallable());
06     }
07     for (int i = 0; i < futures.length; i++) {
08         Integer retVal = futures[i].get();
09         System.out.println(retVal);
10     }
11     executorService.shutdown();
12 }
```

86. Is it possible to perform stream operations in Java 8 with a thread pool?

Collections provide the method

parallelStream()

to create a stream that is processed by a thread pool. Alternatively you can call the intermediate method

parallel()

on a given stream to convert a sequential stream to a parallel counterpart.

87. Is object pooling always a performance improvement for multi-threaded applications?

Object pools that try to avoid the construction of new objects by pooling them can improve the performance of single-threaded applications as the cost for object creation is interchanged by requesting a new object from the pool. In multi-threaded applications such an object pool has to have synchronized access to the pool and the additional costs of lock contention may outweigh the saved costs of the additional construction and garbage collection of the new objects. Hence object pooling may not always improve the overall performance of a multi-threaded application.

Fork/Join

88. How can we access the thread pool that is used by parallel stream operations?

The thread pool used for parallel stream operations can be accessed by



```
ForkJoinPool.commonPool()
```



. This way we can query its level of parallelism with

```
commonPool.getParallelism()
```

. The level cannot be changed at runtime but it can be configured by providing the following JVM parameter:

```
-Djava.util.concurrent.ForkJoinPool.common.parallelism=5
```

.

89. What kind of tasks can be solved by using the Fork/Join framework?

The base class of the Fork/Join Framework

```
java.util.concurrent.ForkJoinPool
```

is basically a thread pool that executes instances of

```
java.util.concurrent.ForkJoinTask
```

. The class

```
ForkJoinTask
```

provides the two methods

```
fork()
```

and

```
join()
```

. While

```
fork()
```

is used to start the asynchronous execution of the task, the method

```
join()
```

is used to await the result of the computation. Hence the Fork/Join framework can be used to implement divide-and-conquer algorithms where a more complex problem is divided into a number of smaller and easier to solve problems.

90. Is it possible to find the smallest number within an array of numbers using the Fork/Join-Framework?

The problem of finding the smallest number within an array of numbers can be solved by using a divide-and-conquer algorithm. The smallest problem that can be solved very easily is an array of two numbers as we can determine the smaller of the two numbers directly by one comparison. Using a divide-and-conquer approach the initial array is divided into two parts of equal length and both parts are provided to two instances of

```
RecursiveTask
```

that extend the class

```
ForkJoinTask
```

. By forking the two tasks they get executed and either solve the problem directly, if their slice of the array has the length two, or they again recursively divide the array into two parts and fork two new RecursiveTasks. Finally each task instance returns its result (either by having it computed directly or by waiting for the two subtasks). The root tasks then returns the smallest number in the array.

91. What is the difference between the two classes

```
RecursiveTask
```

and

```
RecursiveAction
```

?

In contrast to

```
RecursiveTask
```

the method

```
compute()
```

of

```
RecursiveAction
```

does not have to return a value. Hence

```
RecursiveAction
```



can be used when the action works directly on some data structure without having to return the computed value.



[Learn more here:](#)

`java.util.concurrent.RecursiveTask` Example

`java.util.concurrent.ForkJoinPool` Example

`java.util.concurrent.ForkJoinWorkerThread` Example

Concurrent Collections

92. What are Concurrent Collection Classes?

Java Collection classes are fail-fast which means that if the Collection will be changed while some thread is traversing over it using iterator, the `iterator.next()` will throw `ConcurrentModificationException`.

93. What is Java Memory Model?

The Java memory model describes how threads in the Java programming language interact through memory. Together with the description of single-threaded execution of code, the memory model provides the semantics of the Java programming language.

94. What is the difference between

`HashMap`

and

`Hashtable`

particularly with regard to thread-safety?

The methods of

`Hashtable`

are all synchronized. This is not the case for the

`HashMap`

implementation. Hence

`Hashtable`

is thread-safe whereas

`HashMap`

is not thread-safe. For single-threaded applications it is therefore more efficient to use the “newer”

`HashMap`

implementation.

95. Is there a simple way to create a synchronized instance of an arbitrary implementation of

`Collection`

,
`List`

or

`Map`

?

The utility class `Collections` provides the methods

`synchronizedCollection(Collection)`

,
`synchronizedList(List)`

and

`synchronizedMap(Map)`

that return a thread-safe collection/list/map that is backed by the given instance.

[Learn more here:](#)

Java 8 Parallel Arrays Example

`java.util.concurrent.ConcurrentNavigableMap` Example



java.util.concurrent.ConcurrentSkipListMap Example

java.util.concurrent.CopyOnWriteArraySet Example

java.util.concurrent.CopyOnWriteArrayList Example



BlockingQueue

96. What is BlockingQueue?

A BlockingQueue is a java Queue that support operations that wait for the queue to become non-empty when retrieving and removing an element, and wait for space to become available in the queue when adding an element. The interface TransferQueue has been added. It is a refinement of the BlockingQueue interface in which producers can wait for consumers to receive elements.

[Learn more here:](#)

Java BlockingQueue Example

java.util.concurrent.DelayQueue Example

java.util.concurrent.LinkedBlockingQueue Example

Java.util.concurrent.SynchronousQueue Example

java.util.concurrent.ArrayBlockingQueue Example

java.util.concurrent.locks.AbstractQueuedSynchronizer Example

Blocking Queue example to execute commands

Blocking Queue example of limited connection pool

Synchronous Queue example to execute commands

Atomic Variables

97. What do we understand by a CAS operation?

CAS stands for compare-and-swap and means that the processor provides a separate instruction that updates the value of a register only if the provided value is equal to the current value. CAS operations can be used to avoid synchronization as the thread can try to update a value by providing its current value and the new value to the CAS operation. If another thread has meanwhile updated the value, the thread's value is not equal to the current value and the update operation fails. The thread then reads the new value and tries again. That way the necessary synchronization is interchanged by an optimistic spin waiting.

98. Which Java classes use the CAS operation?

The SDK classes in the package

```
java.util.concurrent.atomic
```

like

```
AtomicInteger
```

or

```
AtomicBoolean
```

use internally the CAS operation to implement concurrent incrementation.

```
01 public class CounterAtomic {
02     private AtomicLong counter = new AtomicLong();
03
04     public void increment() {
05         counter.incrementAndGet();
06     }
07
08     public long get() {
09         return counter.get();
10     }
11 }
```

[Learn more here:](#)

java.util.concurrent.atomic.AtomicBoolean Example

java.util.concurrent.atomic.AtomicLongArray Example

Java AtomicIntegerArray Example

Java AtomicInteger Example

Java AtomicMarkableReference Example



Concurrent Random Numbers

99. What is the purpose of the class

java.lang.ThreadLocal

?

As memory is shared between different threads,

ThreadLocal

provides a way to store and retrieve values for each thread separately. Implementations of

ThreadLocal

store and retrieve the values for each thread independently such that when thread A stores the value A1 and thread B stores the value B1 in the same instance of

ThreadLocal

, thread A later on retrieves value A1 from this

ThreadLocal

instance and thread B retrieves value B1.

100. What are possible use cases for

java.lang.ThreadLocal

?

Instances of

ThreadLocal

can be used to transport information throughout the application without the need to pass this from method to method. Examples would be the transportation of security/login information within an instance of

ThreadLocal

such that it is accessible by each method. Another use case would be to transport transaction information or in general objects that should be accessible in all methods without passing them from method to method.

[Learn more here:](#)

java.util.concurrent.ThreadLocalRandom Example

Ok, so now you are ready for your interview! Don't forget to check our FREE Academy course *Java Concurrency Essentials!*

If you enjoyed this, then **subscribe to our newsletter** to enjoy weekly updates and complimentary whitepapers! Also, check out **JCG Academy** for more advanced training!

You are welcome to contribute with your comments and we will include them in the article!

Tagged with:

CONCURRENCY

INTERVIEW

INTERVIEW QUESTIONS

ULTIMATE



(0 rating, 0 votes)

You need to be a registered member to rate this. 8 Comments 6067 Views Tweet it!

Do you want to know how to develop your skillset to become a **Java Rockstar?**

Subscribe to our newsletter to start Rocking right now!
To get you started we give you our best selling eBooks for **FREE!**

1. JPA Mini Book
 2. JVM Troubleshooting Guide
 3. JUnit Tutorial for Unit Testing
 4. Java Annotations Tutorial
 5. Java Interview Questions
 6. Spring Interview Questions
 7. Android UI Design
- and many more

Enter your e-mail...

☐ I agree to the Terms and Privacy Policy

Sign up



LIKE

8

A CODE GEEKS



Join the discussion...



This site uses Akismet to reduce spam. [Learn how your comment data is processed.](#)

☒ Subscribe ▾

▲ newest ▲ oldest ▲ most voted



Guest

luciano



You should include questions about, fail-fast and fail safe iterators, but this post was amazing.

+ 0 - [Reply](#)

🕒 4 years ago



Guest

Vasya



Great questions

+ 0 - [Reply](#)

🕒 4 years ago



Guest

Yusuf



Thanks a lot, it is very beneficial

+ 0 - [Reply](#)

🕒 4 years ago



Guest

Ganesh



Awesome work.

+ 0 - [Reply](#)

🕒 3 years ago



Guest

Abhinav



What is call stack and why we use it with thread ?

+ 0 - [Reply](#)

🕒 3 years ago



Guest

Nadia



Thank you! I have a lot to read! XD I love to learn.

+ 0 - [Reply](#)

🕒 3 years ago





Masudur Rahman

Callable allows the method run() to return a value and to throw an exception. It should be call() to return a value and to throw an exception. It's confusing to me.



+ 0 -

[Reply](#)

🕒 3 years ago

Guest



culey



Thank you for the comprehensive list .Where can I get the pdf?

Guest

+ 0 -

[Reply](#)

🕒 3 years ago

KNOWLEDGE BASE

[Courses](#)

[Examples](#)

[Minibooks](#)

[Resources](#)

[Tutorials](#)

PARTNERS

[Mkyong](#)

THE CODE GEEKS NETWORK

[.NET Code Geeks](#)

[Java Code Geeks](#)

[System Code Geeks](#)

[Web Code Geeks](#)

HALL OF FAME

["Android Full Application Tutorial" series](#)

[11 Online Learning websites that you should check out](#)

[Advantages and Disadvantages of Cloud Computing – Cloud computing pros and cons](#)

[Android Google Maps Tutorial](#)

[Android JSON Parsing with Gson Tutorial](#)

[Android Location Based Services Application – GPS location](#)

[Android Quick Preferences Tutorial](#)

[Difference between Comparator and Comparable in Java](#)

[GWT 2 Spring 3 JPA 2 Hibernate 3.5 Tutorial](#)

[Java Best Practices – Vector vs ArrayList vs HashSet](#)

ABOUT JAVA CODE GEEKS

JCGs (Java Code Geeks) is an independent online community focused on creating the ultimate Java to Java developers resource center; targeted at the technical and non-technical team lead (senior developer), project manager and junior developer. JCGs serve the Java, SOA, Agile and Telecom communities with daily news written by domain experts, articles, tutorials, reviews, announcements, code snippets and source projects.

DISCLAIMER

All trademarks and registered trademarks appearing on Java Code Geeks are the property of their respective owners. Java is a trademark or registered trademark of Oracle Corporation in the United States and other countries. Examples Java Code Geeks is not connected to Oracle Corporation and is not sponsored by Oracle Corporation.

